ENAM Holdings - AI Engineer Build Challenge

FINN: Financial Investment Agent

Release Date: Tuesday, October 27, 2026

Deadline: Sunday, November 1, 2026

Deliverable: A mobile-first working prototype + repository + covering note

What we're evaluating: How you build, what you choose to build, and the creativity and judgment behind it — not completeness or polish.

> **Companion file.** A separate reference document accompanies this brief — *"Reference: Announcement Categories & Financial Glossary."* It lists the corporate-announcement categories and explains every financial term and document type mentioned here in plain English. You do not need any prior finance knowledge — lean on it freely.

About this challenge

We are a public-markets investment firm. We run lean — a small team makes high-conviction decisions and lives with the outcomes. We're building an AI-powered intelligence layer for how we invest, and this challenge is a taste of the real problem. You do not need finance or public-markets expertise: the domain is scaffolded for you, and the companion reference document explains everything financial referenced here. We're hiring for building ability, product instinct, and creativity.

The problem

An investor today is drowning in information — prices, filings, announcements, results, macro. It never stops, and most of it doesn't matter. The scarce skill is not aggregating information. It's filtering it: surfacing the few things that actually matter and burying everything else, beautifully.

Build something that becomes part of a daily investor's workflow — a product they reach for through the day that saves them time and shows them only what's important.

What you'll build

FINN - a mobile-first financial investment agent. It works in two modes of interaction:

- **Push** — it proactively tells you what changed and why you should care.
- **Pull** — you ask it to research or analyse something on demand.

We guide you on the backend (the data). The frontend — the product, the interaction model, the design — is deliberately yours. That's where we most want to see how you think.

Your companies

- **Portfolio:** choose your own mock portfolio of 15–20 NSE/BSE-listed companies. Pick a mix that makes for an interesting demo — different sectors and sizes will show more.
- **Ad-hoc research:** in addition, pick a handful of companies (say 3–5), which can be outside your portfolio, to demo the on-demand research on.

- **Build the engine to generalise.** The research must not be hard-wired to your demo set — build it so it works on any NSE ticker where the data is available. We may point it at a company you didn't pick.

The backend — your data sources (guided)

Three sources. Get them working; don't over-engineer them — this part is plumbing, not the point.

- **Price & volume (end-of-day)** for your portfolio. Freely scrapable at end of day from public sources — Yahoo Finance is the usual starting point (the yfinance Python library; Indian tickers use .NS for NSE and .BO for BSE, e.g. RELIANCE.NS). Anything equivalent works.
- **Corporate announcements / filings** from NSE/BSE for those stocks. A classification of announcement types is provided in the companion reference document — use it, adapt it, or ignore it. It's there to orient you, not to constrain you.
- **Company fundamentals** — annual reports and concall (earnings-call) transcripts, both available as links on screener.in (or the company's investor-relations site).

The frontend — your canvas

This is the heart of the challenge. Give us part of a real investor's daily workflow, built beautifully. The items below are examples of the kinds of things that help — illustrations, not a checklist. What you choose to build, and how, is the signal.

Push — the app reaches out to you

Push is the app telling you what changed without being asked. Think of it in two flavours:

1. A windowed digest. For a time window I choose — by default the last 24 hours, but let me pick the period — show me every corporate filing/announcement across my portfolio, filtered and prioritised so I see what matters first and the routine stuff is collapsed or hidden.

2. Triggered alerts. Tell me the moment something crosses a threshold, for example:

- a material announcement of a type I care about (results, M&A, management change, and so on); or
- an unusual move in my end-of-day data — e.g. volume more than ~2× its recent average, or a price move beyond ±X% in a day. Make the threshold configurable: the point is that I decide what "unusual" means and the app watches for it.

How a pushed item should be presented — two examples we love; the rest is your call:

- **Management change** → don't just say "CFO resigned." Pull a short background/bio of the incoming person (e.g. from LinkedIn) so I instantly know who they are and where they came from.
- **Quarterly results** → take the actual results filing and give me a first cut of the headline numbers — Revenue, EBITDA, PBT and PAT — this quarter vs. last quarter (QoQ) and vs. the same quarter a year ago (YoY). Those four are defined in the reference document.

For every other announcement type, decide what earns my attention: a one-line summary next to the filing, or just the filing surfaced at the right priority. That judgment is the point.

Pull — you ask it to dig in

Pull is on-demand analysis of a company, grounded in its fundamentals (annual reports + concall transcripts). Below are eight kinds of output a good research view might contain — we spell them

out because you may not have a finance background. Treat them as a menu: pick what you build and how deep you go. Each is explained further in the reference document.

- **Business snapshot** — what the company does, its segments, revenue mix.
- **Financial snapshot** — current headline numbers (revenue, margins, PAT, key ratios) at a glance.
- **Trajectory** — the multi-year direction: what's growing, what's compressing.
- **Balance sheet** — debt, leverage, working-capital cycle, funding structure.
- **Cash quality** — cash-flow trend and cash conversion: do profits turn into cash?
- **Narrative vs. numbers** — where management's story and the actual financials agree or diverge.
- **Guidance** — what management has explicitly guided on headline numbers (growth, margins, capex).
- **Bull vs. bear** — a tight case each way.

Worked example — numbers vs. narrative: take the last four quarters and, against management's own headline guidance and growth commentary, tell me what was promised vs. what was actually delivered. This is the difference between summarising management and analysing them — and it's the kind of thing we value most.

One discipline for the pull side: ground it in the primary sources (annual reports + transcripts). No broker notes, no news, no research reports, and don't lean on the model's own outside knowledge of the company. Differentiated signal lives in primary data; secondary commentary is where it gets noisy. (Push-side enrichment like a person's professional background is a different thing and completely fine.)

The interaction model — experiment

How the investor interacts with all of this is entirely up to you, and we strongly encourage you to experiment. Chat, voice, a copilot, a dashboard, or something we haven't thought of — try more than the obvious. Everything mobile-first. The best submissions will show a point of view about how an investor should interact with this, not just what data to show.

The bar — no AI slop

Whatever you build: core, concise, beautifully presented, and interactive. It should save the investor time and show them what's important — and feel like a product someone thought hard about, not a model filling a template. Density and design both count.

Tools, models & cost

- **No restriction on models.** Use anything — open-source, local, or any commercial provider. Pick the right tool for each step; that choice is part of what we're evaluating.
- **This is comfortably doable on free tiers** — you do not need to spend to compete.
- **Efficiency earns bonus points.** The less you spend to achieve the same result, the better. Smart model routing (cheap models for bulk extraction, a stronger model only where reasoning demands it), caching, and chunking are exactly the instincts we're looking for.

Scope, and how we grade

- The deliverable is a mobile-first working prototype. We grade creativity, product sense, judgment, and architecture — **not completeness.**

- **Thin-but-working beats broad-but-broken.** A focused prototype that nails one or two workflows beautifully beats a broad one that doesn't run. A demo that doesn't run is the worst outcome.
- Build the smallest thing that **proves** you can build the real thing well.

What to submit

- **The mobile-first prototype** — a deployed link, Expo/TestFlight, or responsive web; however you ship it, make it easy for us to try on a phone.
- **The GitHub repository**
- **Your AI build history** — the coding-agent and LLM conversations you used while building (e.g. Claude Code, Cursor, or ChatGPT exports / share links). We want to see how you work with the tools, not just the finished code.
- **A short writeup (in your covering email)** Cover what you built, the key decisions and trade-offs, what you deliberately left out, and where you'd take it next.
- Submit your response in an email with the links and attachments to *assessment@enam.com* with the subject line as **"ENAM - IIT screening - 2027"**